DAVID DAMIEN JACINTO

Senior Software Engineer

(432) 220-8122 | jacintoodavid@gmail.com | Waco, TX | LinkedIn

CAREER SUMMARY

Senior Software Engineer with 9+ years building production systems in regulated environments across healthcare AI, enterprise SaaS workflow, insurance operations, and investment banking. Strong in Java, Python, Node.js, Go, and C#, with modern web front ends in React, Angular, and Vue.js. I focus on reliability, auditability, and shipping software that holds up under real production pressure.

TECHNICAL SKILLS

- Languages: Java, JavaScript, TypeScript, Python, C#, SQL, Go
- Backend & APIs: Spring Boot, Spring MVC, Hibernate/JPA, Node.js (Express), REST, gRPC, OAuth2, async workers, idempotency patterns, SQS, Kafka
- Frontend: React, Next.js, Angular, Vue.js, RxJS, WebSockets
- Datastores: PostgreSQL, MySQL, MongoDB, DynamoDB, AWS RDS, Azure SQL, MS SQL Server, Oracle, Redis
- AWS & DevOps: IAM, Lambda, API Gateway, ECS/Fargate, EC2, S3, CloudFormation, CloudWatch, Docker, Kubernetes (EKS), CI/CD
- Search & Observability: Amazon OpenSearch, CloudWatch Logs/Metrics, distributed tracing (OpenTelemetry-style)
- Process & Tools: Git, Agile, Scrum, code review, incident response, runbooks
- AI/ML: RAG pipelines, vector search, LangChain, LangGraph, MLflow, XGBoost, PyTorch, evaluation + grounding checks

PROFESSIONAL EXPERIENCE

SENIOR SOFTWARE ENGINEER | Hippocratic AI | Palo Alto, CA — Jan 2023 – Nov 2025

- Led full-stack delivery across safety architecture, patient education agents, and post-discharge automation, supporting 100,000+ patient conversations per day in production by designing workflows that stayed simple for users and provably safe for clinicians.
- Built core backend microservices in Java with Spring Boot/Spring MVC using gRPC and REST for orchestration, safety checks, and escalation paths, cutting p95 workflow latency by about 30% by tightening call chains and improving caching, timeouts, and retries.
- Developed FastAPI services for inference and orchestration using RAG pipelines and vector search backed by Amazon OpenSearch, improving grounded answer rate by about 15 points on internal evaluation sets by strengthening retrieval ranking, filtering, and grounding checks.
- Implemented multi-agent orchestration patterns with LangChain and LangGraph using explicit state transitions persisted in DynamoDB and PostgreSQL (RDS), improving completion rates by about 20% by making routing deterministic and adding safer tool use and clinician handoffs.
- Built internal React and Next.js tooling in TypeScript with supporting Node.js utilities for configuration, monitoring, and QA, reducing triage time by about 25% with better trace inspection, retrieval-quality views, and guardrail outcome dashboards.
- Designed audit-ready data flows using Hibernate/JPA with PostgreSQL plus S3 for durable artifacts and encrypted logs, improving operability by making every model version, prompt change, and safety decision traceable end-to-end.
- Ran cloud infrastructure on AWS using IAM, API Gateway, Lambda, ECS/Fargate and EC2 with CloudFormation and CloudWatch, improving release reliability and reducing typical rollback time from around 30 minutes to under 10 minutes with stronger health checks and automation.

- Standardized experimentation and lifecycle management with MLflow and compliance-first controls (HIPAA/PHI, encryption, least-privilege access, audit logging, and FHIR-aligned patterns), cutting the experiment-to-production loop from about two weeks to about one week without compromising safety.

SENIOR SOFTWARE ENGINEER | Kizen | Austin, TX — Jun 2021 – Dec 2022

- Delivered full-stack workflow orchestration, commissions automation, and contract lifecycle capabilities using Angular and TypeScript (plus targeted Vue.js modules), supporting dozens of enterprise customers running thousands of workflows per day with reliable audit history.
- Developed backend services in Java with Spring Boot/Spring MVC and C# ASP.NET Core using REST APIs and OAuth2, improving reliability by reducing duplicate-processing incidents by about 40% through idempotency keys, safer retries, and backward-compatible API versioning.
- Built event-driven processing for long-running workflows using queue-based workers and consistent state transitions, improving throughput by about 25% by making replays deterministic and isolating side effects behind durable message handling.
- Implemented data layers across Azure SQL, PostgreSQL, and Redis caching while supporting customer deployments on MySQL, MS SQL Server, and Oracle, cutting high-traffic report queries from seconds to sub-second performance by tuning indexes and query plans.
- Built Angular operations tooling for workflow configuration, RBAC, and permissions simulation, reducing configuration mistakes and support tickets by about 20% by adding validation, audit-history views, and safer defaults.
- Shipped integrations using OAuth2, webhooks, and batch sync jobs with Node.js services where it fit best, improving sync reliability and reducing data-quality issues by about 30% through schema validation, deduplication, and backoff-aware retries.
- Strengthened security and compliance controls for regulated customers by adding encryption, secrets management, and audit logging, supporting SOC 2-aligned practices while keeping developer velocity high with repeatable patterns.
- Improved operability with containerized deployments, Cloud-ready runbooks, and Git-driven CI/CD under Agile/Scrum cadence, reducing incident diagnosis time by centralizing logs/metrics and shipping smaller, safer releases.

BACKEND DEVELOPER | Credit Suisse | Zurich, Switzerland — Aug 2016 – May 2021

- Delivered production ML and data-driven services across markets risk and financial crime, processing millions of order and trade events daily while keeping outputs investigation-friendly with consistent audit trails and clear ownership.
- Built Kafka-based ingestion and enrichment flows in Java and Python, improving data freshness and reducing late-arriving events by about 20% by tightening backpressure handling, retry strategies, and dead-letter patterns.
- Developed Java services in Spring Boot with Hibernate/JPA for scoring and analytics retrieval backed by enterprise relational stores (Oracle and MS SQL Server), keeping p95 scoring latency around 150 milliseconds under load with strong error handling and tracing.
- Built Python feature-engineering pipelines and XGBoost models, reducing false positives by about 15% on a key trade surveillance signal by improving calibration, monitoring drift, and tightening offline-to-online consistency checks.
- Implemented batch and streaming data processing using SQL and distributed compute, cutting a daily risk job from about 4 hours to about 2.5 hours by reworking joins, partitioning, and compute sizing for predictable runtime.
- Created reproducible MLOps workflows with model versioning, automated backtests, and CI checks, reducing release cycle time by about 30% for model updates while keeping governance artifacts complete and reviewable.
- Strengthened governance for regulated finance systems by shipping explainability artifacts (top drivers and reason codes) and access controls, reducing investigator friction by making alerts actionable instead of opaque.
- Partnered with quants, risk managers, and compliance stakeholders in Agile/Scrum teams, translating regulatory expectations into measurable features and dashboards and delivering changes through disciplined Git code review.

EDUCATION

B.S. COMPUTER SCIENCE | TEXAS A&M UNIVERSITY — 2012 – 2016